



Introduction to Machine Learning

Classification

Olatz Arbelaitz: olatz.arbelaitz@ehu.eus
www.aldapa.eus

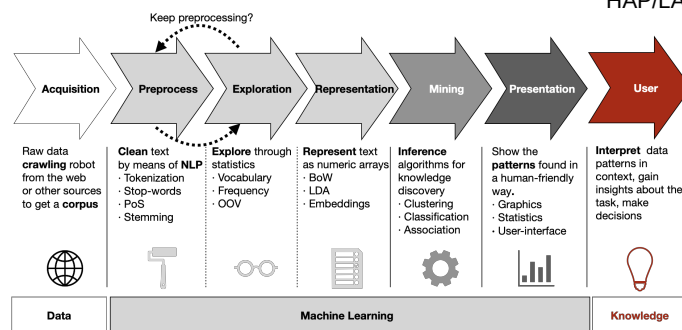


Topics

- 1.- Introduction. Machine Learning for LNP
- 2.- Learning with WEKA software:
 - 2.1.-Introduction
 - 2.2.-Preprocessing
Attribute (feature) selection
 - 2.3.-Evaluation
 - 2.4.- **Basic ML algorithms:** Naive Bayes, K-NN, Decision Trees, Rules, ...

2

Classification



3

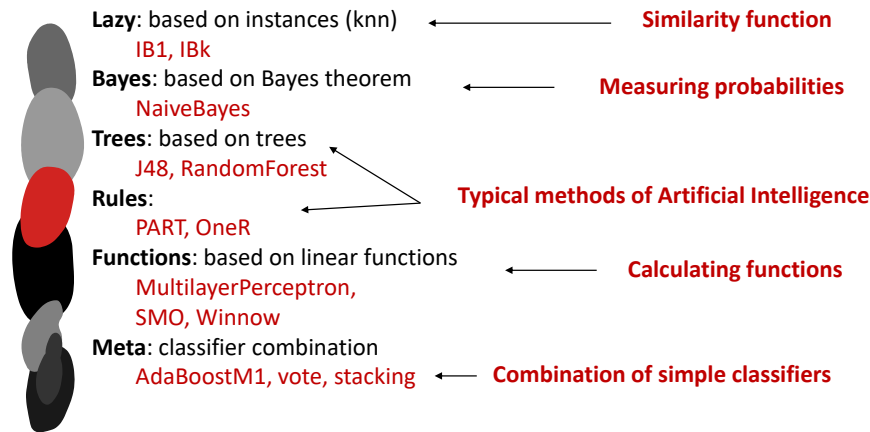
Classification

Classification process

- Division of the corpora (*Test options*)
 - train / test
 - Cross-validation
- **Classifier** (*Classify*)
 - Set parameters
- Evaluation (*Classifier Output*)
 - Confusion matrix
 - Precision/recall
 - Microaveraging/macroaveraging

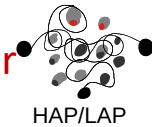
4

Classifiers



5

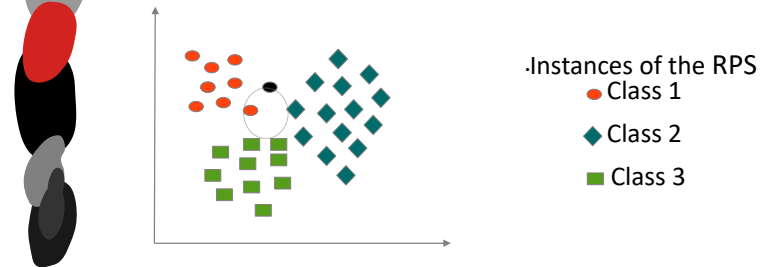
Based on instances: K-NN classifier



k Nearest Neighbour (**k-NN**)

k-NN is a **lazy** classifier. The classifier is based on the learning instances stored in memory, there is **not** model of the categories built

1-NN The class of the new instance to be classified will be the class of its nearest neighbour in the reference pattern set (RPS)



6

Based on instances: K-NN classifier



k Nearest Neighbour (**k-NN**): the majority class within the K nearest instances in the RPS is chosen

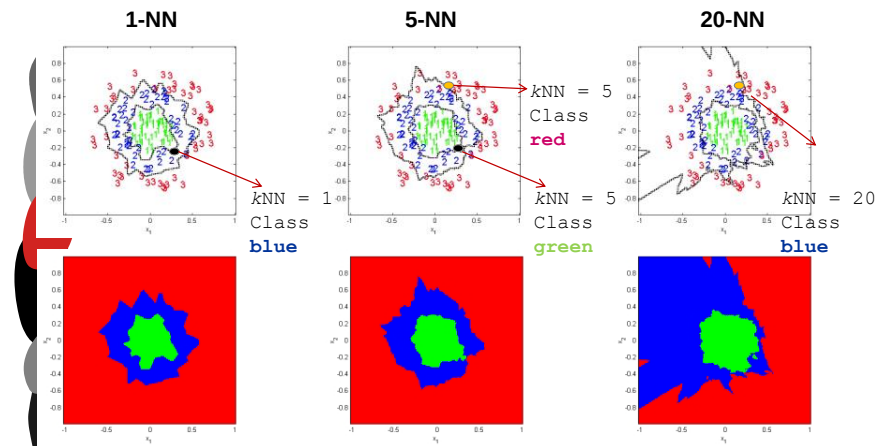
Procedure to classify new instances

- 1.- Calculate the **distance** between the instance to classify and all the instances in the RPS
For example Euclidean distance (n dimensions)
- 2.- Select the **k nearest** instances (smallest distance)
- 3.- Assign as class the **majority class** within the k instances

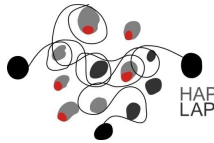
Main parameters

- Value of **k**, **number of examples** used to classify
- Distance** or similarity measure used to compare instances
- Criteria to select** the k nearest instances
- Criteria to decide** the class of the new instances

kNN versus 1NN



8



Based on instances

Weka: by default distance: **Euclidean**

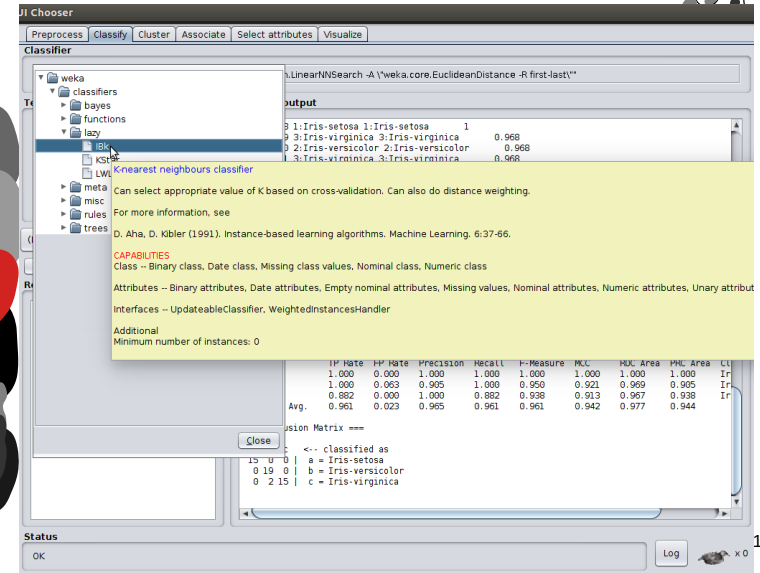
measures how far/near the elements are in a vector space

In **Lazy** classifiers:

- **IBk**: the k most similar. Normalizes numeric attributes (between 0-1) to apply distances. If $K > 1$ we can weight distances.
 - Many distance options
- **windowSize**: can be used to limit the number of instances to be kept

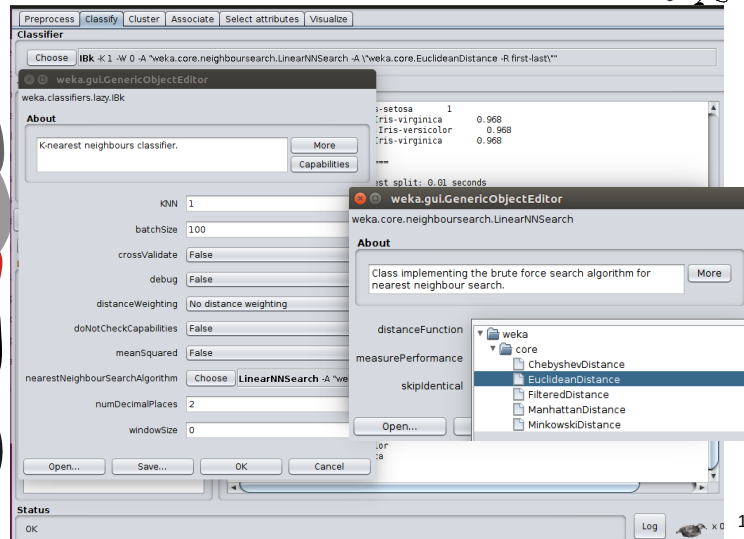
9

Classifier



10

Classifier



11

k-NN. Example

Example (text classification)

category

- d₁**: ill(3), player, doctor(5), health(2) **health**
- d₂**: nurse, play(3), doctor(4), health(2) **health**
- d₃**: player(4), play(2), doctor(2), ball(3) **sport**
- d₄**: ill(3), play(4), doctor(2), health **???**

dictionary

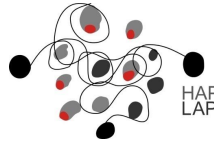
Nurse
Ill
Player
Play
Doctor
Health
Ball

$$|d_j, d_z| = \sqrt{\sum_{i=1}^n (w_{ji} - w_{zi})^2}$$

- d₁**: 0 3 1 0 5 2 0 |d₁-d₁| =
- d₂**: 1 0 0 3 4 2 0 |d₂-d₁| =
- d₃**: 0 0 4 2 2 0 3 |d₃-d₁| =
- d₄**: 0 3 0 4 2 1 0

12

Knn. Assignment



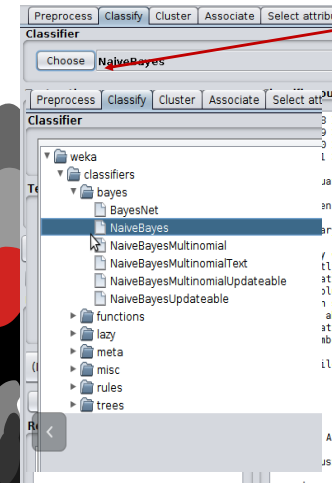
Open ReutersGrainTrain

After a feature selection process, for the best set try k-NN algorithm and adjust K parameter

| | k = 1 | k = 3 | k = 5 | K = |
|------------------------|-------|-------|-------|-----|
| RGT_(percentage_split) | | | | |
| RGT_(CV-10 fold) | | | | |

14

Types of classifiers



Lazy: based on instances (knn)

IB1, IBk

Bayes: based on Bayes theorem

NaiveBayes

Trees: based on trees

J48, NBTree, RandomForest

Rules:

PART, OneR

Functions: based on linear functions

MultilayerPerceptron,

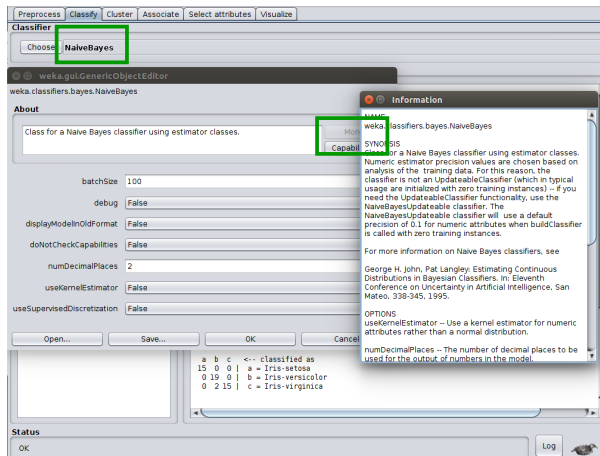
SMO, Winnow

Meta: classifier combination

AdaBoostM1, vote, stacking

15

Naive Bayes

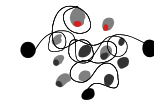


To see the concrete Parameters of a classifier and a short description about them click in the name of the classifier in bold and then click the "More" button

Example: document classification(spam)

http://en.wikipedia.org/wiki/Naive_Bayes_classifier#Document_Classification

Naive Bayes



Classifier based on Bayes Theorem making calculations simpler when:

The database has many features

There are not enough examples to calculate probabilities of all feature combinations.

Revision:

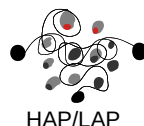
P(wi): prior probability of class wi
(quantifies the probability of a class without any extra information)

P(x|wi): density function probability conditioned to the class
(quantifies the probability of x having a concrete value knowing the class it belongs to)

P(wi|x): posterior probability
(quantifies the probability of an instance to belong to a class)

p(x): probability of instance x (unconditional)
(distribution of the instances)

Naive Bayes



Bayes theorem

$$0 \leq P(w_i) \leq 1 \quad \sum_{i=1, \dots, c} P(w_i) = 1 \quad \sum_{i=1, \dots, c} P(x | w_i) = 1$$

Given the prior probabilities and the density functions the **posterior probability** can be calculated (**≈classification**):

$$P(w_i | x) = \frac{P(x | w_i) * P(w_i)}{P(x)}$$

where

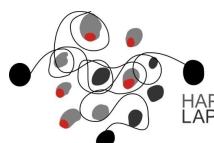
$$P(x) = \sum_{i=1}^c P(x | w_i) * P(w_i)$$

Classification with Naive Bayes

$$w_{NB} = \arg \max_{w_i \in C} P(w_i) \prod_{k=1, \dots, F} P(x_k | w_i)$$

$$\hat{y} = \arg \max_{k \in \{1, \dots, k\}} p(class_k) \prod_{i=1} p(x_i | class_k)$$

NaiveBayes



Bayes theorem: calculating conditional probabilities

$$P(A|B) = P(A) \times P(B|A) / P(B)$$

P(A): Prob. of A, P(A|B): Prob. of A given B is true and P(B|A): Prob. of B given A is true.

Example: is the web page containing word **mode** in French or in English?

And the one containing **maison**?

$$P(\text{French}|\text{mode}) = P(\text{French}) \times P(\text{mode}|\text{French}) / P(\text{mode}).$$

Given a word (*mode*), the probability of the document being in *French* **P(French|mode)** is the following: probability of the document to be in *French* **P(French)** and the probability of having word *mode* if the document is in *French* **P(mode|French)** divided by the proportion of documents containing word *mode* **P(mode)**.

$$P(\text{French}) = 0.08 \quad P(\text{mode}|\text{French}) = 0.62 \quad P(\text{mode}) = 0.15 \quad | \quad P(\text{maison}|\text{French}) = 0.92 \quad P(\text{maison}) = 0.08$$

$$P(\text{French}|\text{mode}) = 0.08 \times 0.62 / 0.15 = 0.33$$

$$P(\text{French}|\text{maison}) = 0.08 \times 0.92 / 0.08 = 0.92$$

20

Naive Bayes

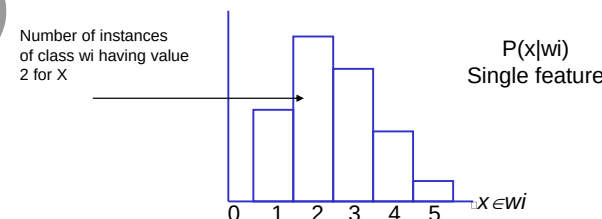


Features are supposed to be **independent**

As a consequence the probability for many features can be calculated as follows

$$P(x_1, \dots, x_F | w_i) = P(x_1 | w_i) \cdot \dots \cdot P(x_F | w_i)$$

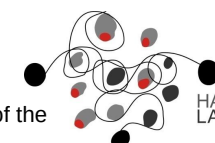
If *x_f* is **discrete** **P(x_f | w_i)** is estimated based on the relative instance frequency of class *w_i* taking value *x_f* (mass function, histograms)



If *x_f* is **continuous**

1. Discretize and treat it as discrete
2. estimate **P(x_f | w_i)** assuming a gaussian (normal) distribution- (only mean and variance required)

NaiveBayes



- **Bag of Words assumption** => assume the position of the words in the document doesn't matter.
- **Conditional Independence** => Assume the feature probabilities **P(x_i | c_j)** are independent given the class **c**.

- To calculate **P(d | c) x P(c)**, we calculate **P(x_i | c)** for each *x_i* in **d**, and multiply them together.
- Then we multiply the result by **P(c)** for the current class. We do this for each of our classes, and choose the class that has the maximum overall value.

$$P(c_i) = [\text{N documents that have been classified as } c_i] / [\text{N documents}]$$

$$P(w_i | c_j) = [\text{count}(w_i, c_j)] / [\sum_{w \in V} \text{count}(w, c_j)]$$

Laplace Smoothing

adding 1 to the numerator and modifying the denominator as such:

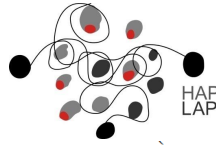
$$P(w_i | c_j) = [\text{count}(w_i, c_j) + 1] / [\sum_{w \in V} (\text{count}(w, c_j) + 1)]$$

$$P(w_i | c_j) = [\text{count}(w_i, c_j) + 1] / [\sum_{w \in V} (\text{count}(w, c_j)) + |V|]$$

where **|V|** is our vocabulary size (

21

NaiveBayes



| | Doc | Words | Class |
|----------|-----|-------------------------------------|-------|
| Training | 1 | Chinese Beijing Chinese | c |
| | 2 | Chinese Chinese Shanghai | c |
| | 3 | Chinese Macao | c |
| | 4 | Tokyo Japan Chinese | j |
| Test | 5 | Chinese Chinese Chinese Tokyo Japan | ? |

Priors:

$$P(c) = \frac{3}{4} \quad \frac{1}{4}$$

$$P(j) = \frac{1}{4}$$

Conditional Probabilities:

$$P(\text{Chinese} | c) = (5+1) / (8+6) = 6/14 = 3/7$$

$$P(\text{Tokyo} | c) = (0+1) / (8+6) = 1/14$$

$$P(\text{Japan} | c) = (0+1) / (8+6) = 1/14$$

$$P(\text{Chinese} | j) = (1+1) / (3+6) = 2/9$$

$$P(\text{Tokyo} | j) = (1+1) / (3+6) = 2/9$$

$$P(\text{Japan} | j) = (1+1) / (3+6) = 2/9$$

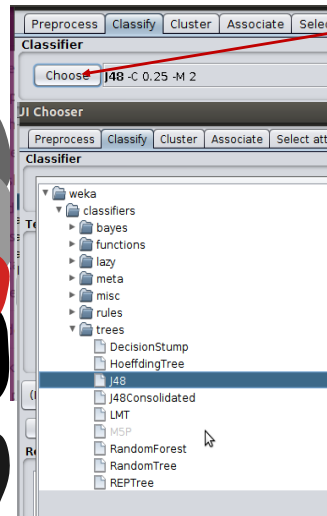
Choosing a class:

$$P(c | d5) \propto 3/4 * (3/7)^3 * 1/14 * 1/14 \approx 0.0003$$

$$P(j | d5) \propto 1/4 * (2/9)^3 * 2/9 * 2/9 \approx 0.0001$$

22

Types of classifiers



Lazy: based on instances (knn)

IB1, IBk

Bayes: based on Bayes theorem

NaiveBayes

Trees: based on trees

J48, NBTre, RandomForest

Rules:

PART, OneR

Functions: based on linear functions

MultilayerPerceptron,
SMO, Winnow

Meta: classifier combination

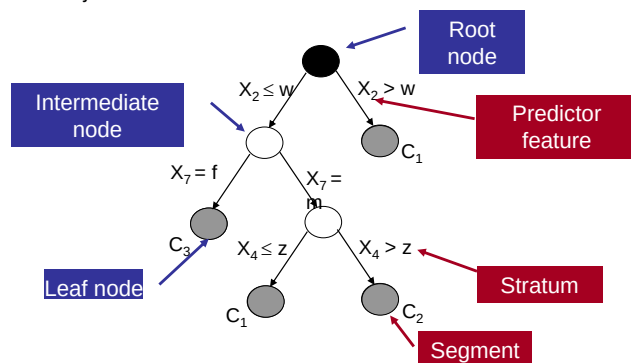
AdaBoostM1, vote, stacking

23

Decision trees



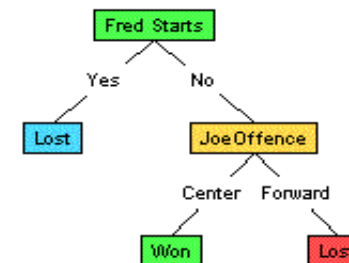
- Based on "Divide and conquer" algorithm
- A classifier in the form of a tree structure
 - Decision node: specifies a test on a single attribute
 - Leaf node: indicates the value of the target attribute
 - Arc/edge: split of one attribute
 - Path: a disjunction of test to make the final decision



Decision trees



- Decision trees classify instances or examples by starting at the root of the tree and moving through it until a leaf node.



- Problem: **Overfitting** (good in learning, worse in generalization)

Decision trees



Input: Training set E (labelled instances)

Output: Decision tree (T)

Algorithm

begin

If all the examples in E are of the same category C_j

then Result **simple node** labelled as C_j

else

begin

Select a **feature** X_i with values x_{i1}, \dots, x_{il}

Partition E in E_1, \dots, E_l according to the values of X_i

Build **subtrees** T_1, \dots, T_l for E_1, \dots, E_l

The result is a **tree with root X_i and subtrees T_1, \dots, T_l**

The branches between X_i and the subtrees are labelled with x_{i1}, \dots, x_{il}

end

end

Decision trees



- **At each node:** selection of an attribute to split- choosing the most “useful” attribute for classifying examples.
 - Nominal features: as many branches as values
 - Numeric features: $\leq, > // >, =, < // <, \text{segment}, >$
- How to decide what is “useful”? Example: information gain
 - measures how well a given attribute separates the training examples according to their target classification
 - At each node, choose to divide the attribute with the largest information gain
- Stopping rule
 - Every attribute has already been included along this path through the tree, or
 - The training examples associated with this leaf node all have the same target attribute value (i.e., their entropy is zero).

Classifier

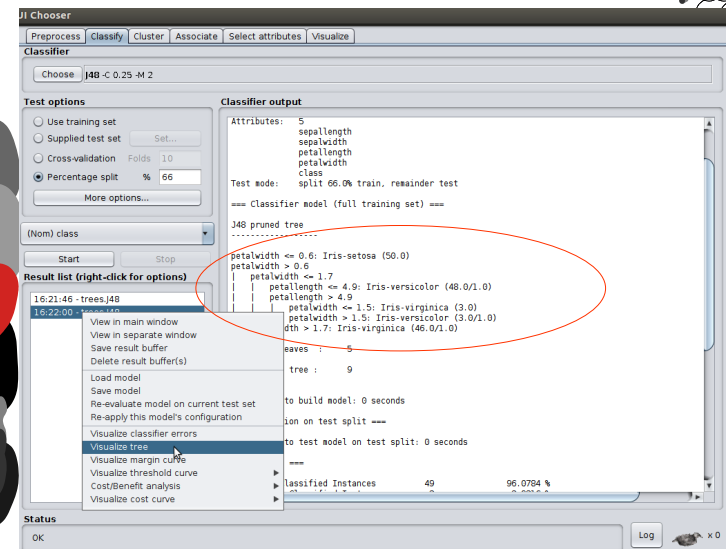
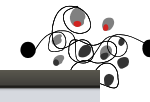


Implementation of C4.5 (Quilan): J48

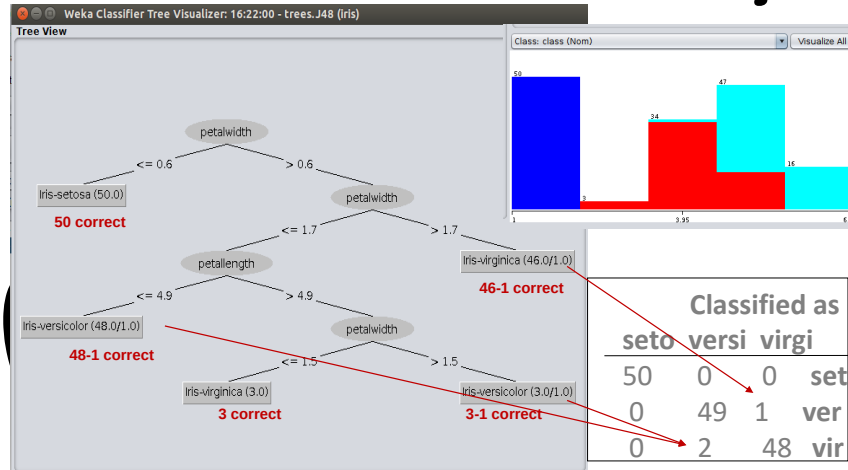
Two branches in numeric attributes ($\leq, >$) see *iris.arff*

With nominal attributes: every value see *soybean.arff*

Decision trees: J48

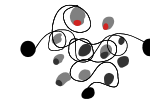


Visualize tree



30

J48 soybean.arff



leafspot-size = lt-1/8

canker-lesion = dna

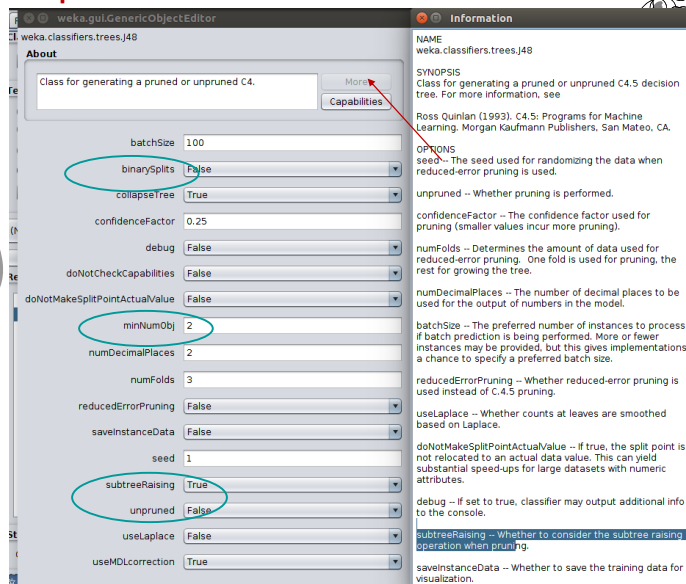
leafspots-marg = w-s-marg
 seed-size = norm: bacterial-blight (21.0/1.0)
 seed-size = lt-norm: bacterial-pustule (3.23/1.23)
 leafspots-marg = no-w-s-marg: bacterial-pustule (17.91/0.91)
 leafspots-marg = dna: bacterial-blight (0.0)
 canker-lesion = brown: bacterial-blight
 canker-lesion = dk-brown-blk: phytophthora-rot (4.78/0.1)
 canker-lesion = tan: purple-seed-stain (11.23/0.23)

....

| Selected attribute | | | |
|---------------------|--------------|----------------|---------------|
| Name: canker-lesion | | Distinct: 4 | Type: Nominal |
| Missing: 38 (6%) | | Unique: 0 (0%) | |
| No. | Label | Count | Weight |
| 1 | dna | 320 | 320.0 |
| 2 | brown | 83 | 83.0 |
| 3 | dk-brown-blk | 177 | 177.0 |
| 4 | tan | 65 | 65.0 |

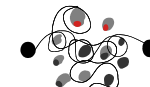
31

J48. Options



32

J48. Options



Pruning

Prepruning (forward pruning)

decide where to cut when building the tree

Postpruning (backward pruning)

- generate the tree and then analyze to decide where to cut

- most used option

- Two options in each node

subtree replacement → replace with leaves

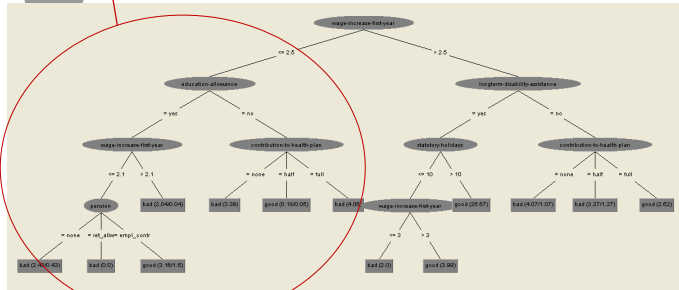
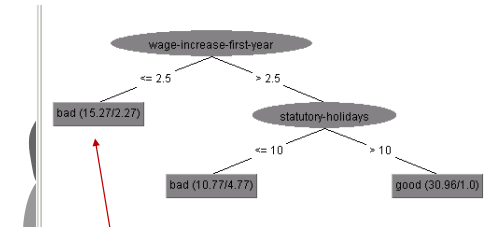
subtree raising → remove subtrees and replace with the lower part → reclassify → time

33

J48. Options. *labor.arff*

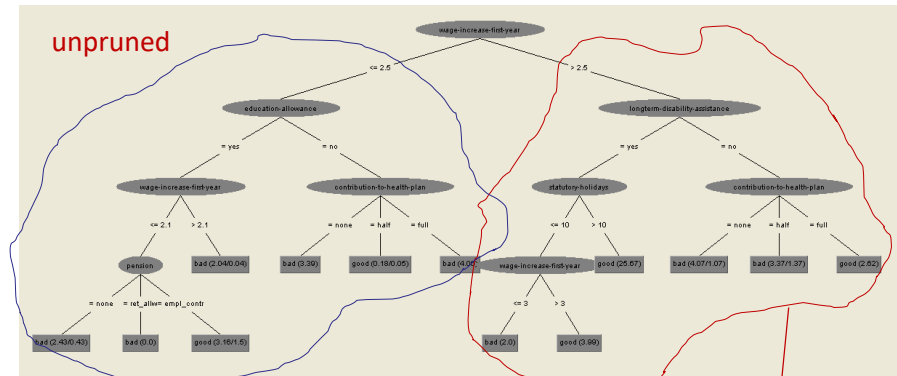


Subtree replacement
5 leaves bad → bad
2 leaves good
F-measure = 0.89



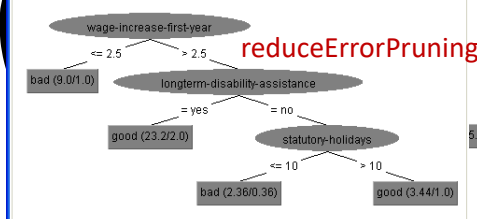
unpruned
F-measure = 0.89

34

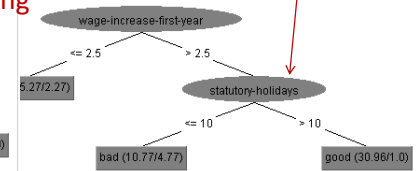


unpruned

Subtree Raising



reduceErrorPruning

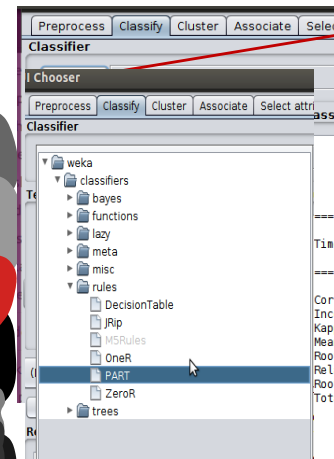
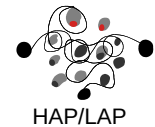


Other tree options



- **Random Tree:**
 - K features are selected randomly in each node (try *soybean.arff*)
- **ADTree:** Freund, Y., Mason, L (1999)
 - Two category branches
 - Boosting iterations: adds 3 nodes in each iteration
- **NBTrees:** Ron Kohavi (1996)
 - NB classifiers in leaf nodes
- **DecisionStump:**
 - Generates binary trees of a single level.
 - To use in Boosting methods
- **Id3:** R. Quinlan (1986)
 - Only nominal attributes
 - Information gain → Gain Ratio
- **Random Forest:**
 - forest of random trees
 - Bagging (multiple classifier system)

Types of classifiers



Lazy: based on instances (knn)

IB1, IBk

Bayes: based on Bayes theorem

NaiveBayes

Trees: based on trees

J48, NBTree, RandomForest..

Rules:

PART, OneR....

Functions: based on linear functions

MultilayerPerceptron,

SMO, Winnow

Meta: classifier combination

AdaBoostM1, vote, stacking

36

37

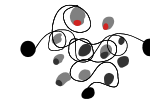
Decision rules



- Divide-and-conquer technique
- Find rules that group instances of a class and discard those that are not of the class
- Overfitting (good in learning but not generalization capacity)
- Maximize the ratio: p/t
 t : number of examples covered by the rule and p those that are positive among them
- Information gain: $p[\log p/t - \log P/T]$
 gain with the new rule
 t and p the same, and P and T same value after introducing the new rule
- To introduce new rules:
 - as many positive examples as possible
 - as few negative as possible

38

Decision rules



- Rule induction

Example iris: 3 rules

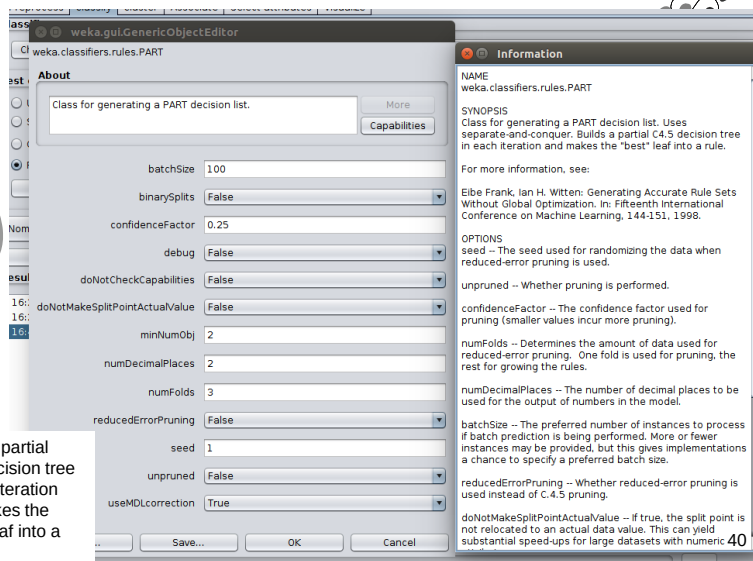
```
if petalwidth <= 0.6 then Iris-setosa
else if petalwidth <= 1.7 AND petallength <= 4.9
then Iris-versicolor
else Iris-virginica
```

Example TC

```
cyclist & Sky & ... then sport
else if crisis & euro & ... then economy
else ...
```

39

Decision rules: PART



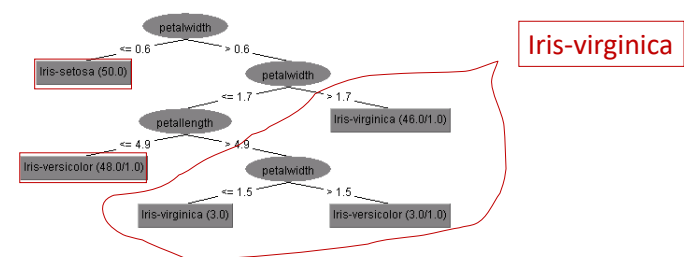
Builds a partial C4.5 decision tree in each iteration and makes the "best" leaf into a rule

Decision rules: PART



HAP/LAP

PART decision list----- 3 rules
 petalwidth <= 0.6: Iris-setosa (50.0)
 petalwidth <= 1.7 AND petallength <= 4.9: Iris-versicolor (48.0/1.0)
 : Iris-virginica (52.0/3.0)



Iris-virginica

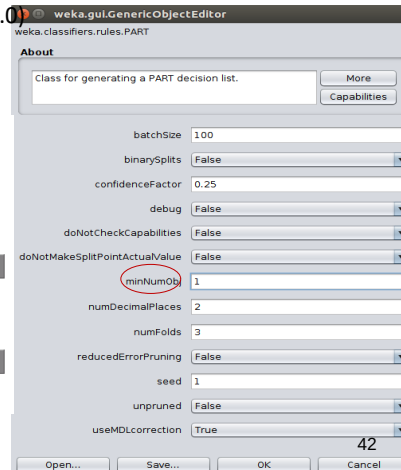
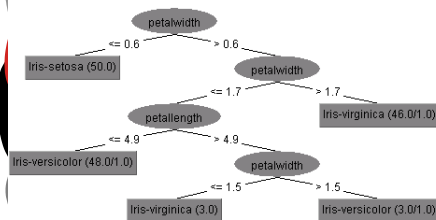
41

Decision rules: PART



HAP/LAP

petalwidth <= 0.6: Iris-setosa (50.0)
 petalwidth > 1.7 AND petallength > 4.8: Iris-virginica (43.0)
 petallength <= 4.7 AND petalwidth <= 1.5: Iris-versicolor (42.0)
 sepalwidth <= 3: Iris-virginica (11.0/4.0)
 : Iris-versicolor (4.0)
 Number of Rules : 5



Decision rules: other options



HAP/LAP

JRIP: similar to RIPPER (Repeated Incremental Pruning to Produce Error Reduction), in accuracy, number of rules and time. Not in memory

OneR: R.C. Holte (1993)

Generates a single rule

Prediction based on the feature with minimum error

Discretizes numeric attributes

Simple rules obtain often better results

DecisionTable: Ron Kohavi (1995)

Uses Best-first search to explore feature set

Cross-validation for evaluation

1-nn can be used when instances can
 not be classified with the information in
 the table

Rules:

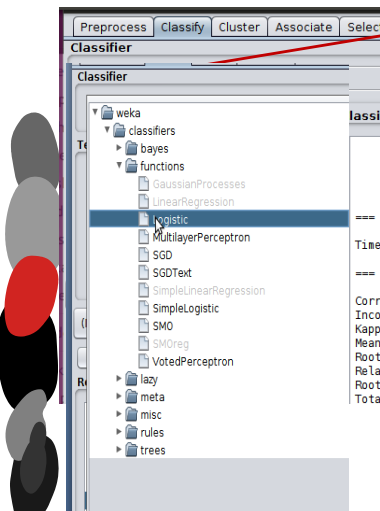
```
petalwidth  class
=====
'(1.75-inf)' Iris-virginica
'(0.8-1.75)' Iris-versicolor
'(-inf-0.8]' Iris-setosa
=====
```

43

Types of classifiers



HAP/LAP



Lazy: based on instances (knn)

IB1, IBk

Bayes: based on Bayes theorem

NaiveBayes

Trees: based on trees

J48, NBTree, RandomForest..

Rules:

PART, OneR....

Functions: based on linear functions

MultilayerPerceptron,

SMO, Logistic

Meta: classifier combination

AdaBoostM1, vote, stacking

Functions



HAP/LAP

Classifiers that can be represented with mathematical equations:

Regression

To predict based on numeric classes and attributes

Ex.: cost of a flat based on size, number of bedrooms, ...

| Size (feet ²) | Number of bedrooms | Number of floors | Age of home (years) | Price (\$1000) |
|---------------------------|--------------------|------------------|---------------------|----------------|
| 2104 | 5 | 1 | 45 | 460 |
| 1416 | 3 | 2 | 40 | 232 |
| 1534 | 3 | 2 | 30 | 315 |
| 852 | 2 | 1 | 36 | 178 |
| ... | ... | ... | ... | ... |

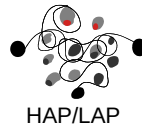
Classification

The values of the class are discrete

44

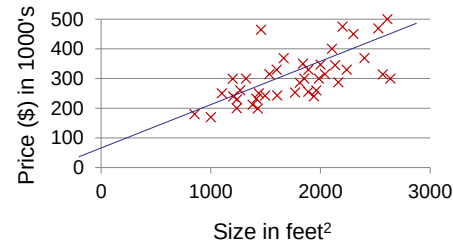
45

Functions



Example: cost of the flat according to size
A single attribute: size(x)

| Size in feet ² (x) | Price (\$) in 1000's (y) |
|-------------------------------|--------------------------|
| 2104 | 460 |
| 1416 | 232 |
| 1534 | 315 |
| 852 | 178 |
| ... | ... |



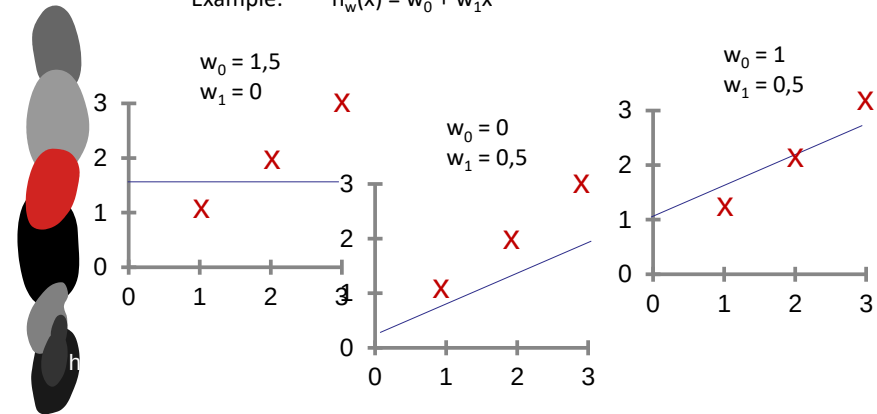
Linear function (hypothesis)
 $h_w(x) = w_0 + w_1x$ (w_i = weights)

46

Functions

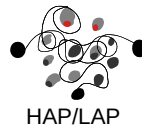


Example: $h_w(x) = w_0 + w_1x$

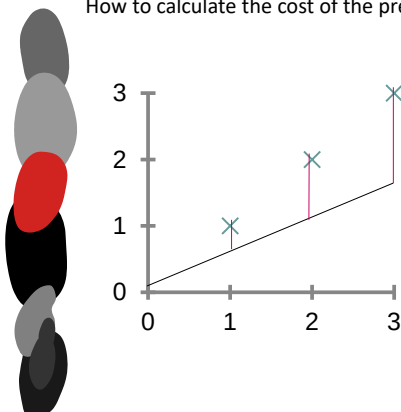


47

Functions



How to calculate the cost of the prediction?



Linear function
 $h_w(x) = w_0 + w_1x_1$
 w_i = weights

Cost function:
 $J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2$

m: number of instances

48

Functions



Linear regression

To predict with numeric classes and attributes.

- **The class** is represented as a linear combination of the features with predicted weights

$$C = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

weights are calculated from the training set

class weights features

- **The predicted class** for each example is calculated in the following way: ($x_0 = 1$)

$$w_0x_0^{(1)} + w_1x_1^{(1)} + w_2x_2^{(1)} + \dots + w_nx_n^{(1)} = \sum w_jx_j^{(1)}$$

find w_j coefficients that minimize the squared difference between the predicted value and the real value

$$\sum_{i=1}^m (k^{(i)} - \sum_{j=0}^n w_j a_j^{(i)})^2$$

Predicted value

m: instances
n: number of features

Real value

49

Example: CPU.arff



Data

| MYCT | MMIN | MMAx | CACH | CHMIN | CHMAX | class |
|------|-------|--------|------|-------|-------|-------|
| 125, | 256, | 6000, | 256, | 16, | 128, | 198 |
| 29, | 8000, | 32000, | 32, | 8, | 32, | 269 |
| 29, | 8000, | 32000, | 32, | 8, | 32, | 220 |
| 29, | 8000, | 32000, | 32, | 8, | 32, | 172 |

....

To calculate w coefficients, the following values need to be minimized:

$$(198 - \sum w_j x_j)^2 + (269 - \sum w_j x_j)^2 + (220 - \sum w_j x_j)^2 + (172 - \sum w_j x_j)^2 + \dots$$

```
class =
0.0491 * MYCT +
0.0152 * MMIN +
0.0056 * MMAx +
0.6298 * CACH +
1.4599 * CHMAX +
-56.075
```

50

Functions: Artificial Neural Networks



Idea:

- model mathematically the human intellectual capacities.
- massively parallel computation schemas
- Unstable classifiers appropriate for multiple classifier systems

Universal approach property:

Some of the models, (Multilayer Perceptron (MLP), Radial Basis Function (RBF)), if an infinite number of patterns can be used, have the capacity to approach any discriminate function with a certain precision.

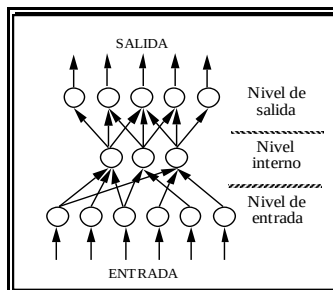
51

Functions: Artificial Neural Networks



Multilayer Perceptron (MLP)

Is a **feedforward** network where every neuron in a layer is connected to every neuron in next layer. The structure would be the following:



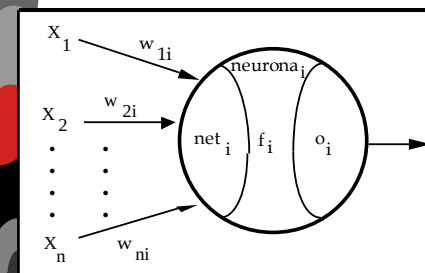
52

Functions: Artificial Neural Networks



The basic structure of an ANN is a single neuron

When the ANN has a single neuron it is called **simple linear perceptron**.

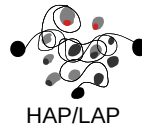


$$o = f\left(\sum_{j=1}^N w_j * x_j + w_0\right)$$

f is an identity type function, sign, sigmoid, etc.

53

Functions



Perceptron

At the beginning, all the weights of the features of category c_i are identical: w_{ki}

For new examples to learn (d_j) the classifier classifies with the weights the classifier has and:

If the classification is correct: no action

If the classification is not correct :

If $d_j \in c_i$ then $w_{ki} := w_{ki} + \alpha$ ($\alpha > 0$)

If $d_j \notin c_i$ then $w_{ki} := w_{ki} - \alpha$ ($\alpha > 0$)

($w_{ki} \rightarrow$ all t_k where $w_{kj} = 1$)

If the weight of the feature (w_{ki}) diminishes in the learning process, feature (t_k) is not useful for classification and it can be removed (on-the-fly term space reduction)

54

Perceptron. Example



- Document classification
 - $d_1 = \dots$ book presentation in Koldo Mitxelena ...
 - $d_1 = \dots$ Europe's economy... the euro has risen
 - $d_j = \dots$ the writer will talk about the book in Koldo Mitxelena
- Representation (lemmas)
 - $d_1 = \dots$ book present Koldo Mitxelena ...
 - $d_1 = \dots$ Europe economy euro have rise...
 - $d_j = \dots$ writer talk about book Koldo Mitxelena ...
- Features:
 - {writer, book, present, novel, ... , read, Koldo, Mitxelena, have, Europe, economy, euro, rise, ...}**
- Category:
 - Culture**

55

Perceptron. Example



- $\{x_1=\text{writer}, x_2=\text{book}, x_3=\text{present}, x_4=\text{novel}, x_5=\text{read}, x_6=\text{Koldo}, x_7=\text{Mitxelena}, x_8=\text{have}, x_9=\text{Europe}, x_{10}=\text{economy}, x_{11}=\text{euro}, x_{12}=\text{rise}, \dots\}$

$d_i = (x_1, x_2, \dots, x_{12}) = \{0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0\}$ **Culture**

$d_i = \{0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1\}$ **Economy**

$d_j = \{1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0\}$ **Culture**

- Algorithm

Beginning $w_k = 0.1$; $x_0 = 1$; $\alpha = 0.8$

$d_i \rightarrow f(x) = x_0 w_0 + x_1 w_1 + x_2 w_2 + \dots + x_{12} w_{12} = 0.5$ $f(x) > 0 \rightarrow$ culture YES

$d_i \rightarrow f(x) = x_0 w_0 + x_1 w_1 + x_2 w_2 + \dots + x_{12} w_{12} = 0.6$ $f(x) > 0 \rightarrow$ **culture YES**

Error \rightarrow recalculating weights ($-\alpha$) $w_0 - w_7 = 0.1$; $w_8 - w_{12} = -0.7$

$d_j \rightarrow f(x) = x_0 w_0 + x_1 w_1 + x_2 w_2 + \dots + x_{12} w_{12} = -0.2$ $f(x) < 0 \rightarrow$ **culture NO**

Error \rightarrow recalculating weights ($+\alpha$) $w_0 - w_2 = 0.9$; $w_3 - w_5 = 0.1$;

$w_6 - w_7 = 0.9$; $w_8 = 0.1$; $w_9 - w_{12} = -0.7$

- Classifier

hyperplane \rightarrow weight vector

56

Functions



Winnow

To recalculate weights ($\alpha > 1, 0 < \beta < 1$)

Multiplication instead of addition subtraction

Positive winnow

If the classification is correct: no action

If the classification is not correct

If $d_j \in c_i$ then $w_{ki} := w_{ki} \times \alpha$ ($\alpha > 1$)

If $d_j \notin c_i$ then $w_{ki} := w_{ki} \times \beta$ ($0 < \beta < 1$)

Balanced winnow

Two weights for each term (+ and -)

If the classification is not correct:

if $d_j \in c_i$ then $w_{ki}^+ := w_{ki}^+ \times \alpha$ ($\alpha > 1$)

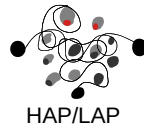
$w_{ki}^- := w_{ki}^- \times \beta$ ($0 < \beta < 1$)

if $d_j \notin c_i$ then $w_{ki}^+ := w_{ki}^+ \times \alpha$

$w_{ki}^- := w_{ki}^- \times \beta$

57

Functions



SVM (Support Vector Machine)

Problems of linear classifiers:

- The boundaries between two classes are linear
- Too simple for many practical applications

SVM- uses linear models to define non linear boundaries between classes

- Projects the input data using non linear mapping
- Converts the instance space to another space
- In the new space, the classes are linearly separable but in the original they are not.

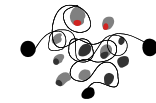
Linear model: *maximun margin hyperplane*

Weka: *functions/SMO classifier* it is slow, CacheSize = 0

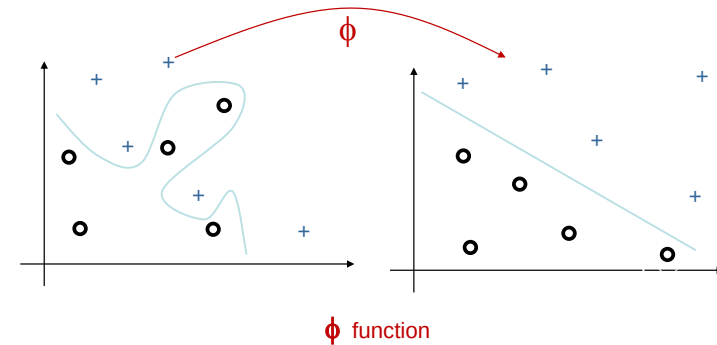
Call weka with: `java -Xms512M -jar weka.jar`

58

Functions

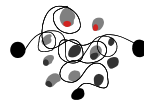


- Hyperplane discriminators: separate positive and negative examples by an hyperplane



59

Kernels



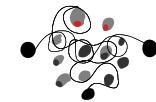
- **Kernel** Function defined in the new space

$$\text{any } x, z \in X \quad K = \langle \phi(x) \cdot \phi(z) \rangle$$

- The Kernel represents the similarity between two objects
- Depending on function f different kernels can be generated different learning algorithms can be used
- Advantages:
 - Not necessary to maintain the feature vectors of the data
 - Adequate to work with complex spaces
 - Good result with high dimensional problems(TC)
 - Efficient ways of calculating internal product exist

60

Kernels



- Document classification
 - $d_i = \dots$ book presentation in Koldo Mitxelena ...
 - $d_i = \dots$ Europes' economy... the euro has risen
 - $d_j = \dots$ the writer will talk about the book in Koldo Mitxelena
- Representation (lemas)
 - $d_i = \dots$ book present Koldo Mitxelena ...
 - $d_i = \dots$ Europe economy euro have rise...
 - $d_j = \dots$ writer talk about book Koldo Mitxelena ...
- Features:
 - {**writer, book, present, nobel, ... , read, Koldo, Mitxelena, have, Europe, economy, euro, rise, ...**}

61

Kernels



$d_1 = x \rightarrow \phi(x) = \{0, 1, 1, 0, 0, \dots, 0, 1, 1, 0, 0, 0, 0\}$
 $d_2 = z \rightarrow \phi(z) = \{1, 1, 0, 0, 0, \dots, 0, 1, 1, 1, 0, 0, 0\}$
 $d_3 = v \rightarrow \phi(x) = \{0, 0, 0, 0, 0, \dots, 0, 0, 0, 1, 1, 1, 1\}$

Kernel function to measure similarity between d_1 and d_2

$d_1 = x \rightarrow \phi(x) = \{0, 1, 1, 0, 0, \dots, 0, 1, 1, 0, 0, 0, 0\}$
 $d_2 = z \rightarrow \phi(z) = \{1, 1, 0, 0, 0, \dots, 0, 1, 1, 1, 0, 0, 0\}$

$$k(x, z) = \langle \phi(x) \cdot \phi(z) \rangle = 3 \text{ identical words}$$

Kernel function to measure similarity between d_2 and d_3

$d_2 = z \rightarrow \phi(z) = \{1, 1, 0, 0, 0, \dots, 0, 1, 1, 1, 0, 0, 0\}$
 $d_3 = x \rightarrow \phi(x) = \{0, 0, 0, 0, 0, \dots, 0, 0, 0, 1, 1, 1, 1\}$

$$k(x, z) = \langle \phi(x) \cdot \phi(z) \rangle = 1 \rightarrow \text{a single word}$$

Kernels



- In the new space of the learning set, the **similarity** of each document with the rest of documents is represented.

$$K_{\text{training}} = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & k(\mathbf{x}_1, \mathbf{x}_2) & \dots & k(\mathbf{x}_1, \mathbf{x}_m) \\ k(\mathbf{x}_2, \mathbf{x}_1) & k(\mathbf{x}_2, \mathbf{x}_2) & \dots & k(\mathbf{x}_2, \mathbf{x}_m) \\ \dots & \dots & \dots & \dots \\ k(\mathbf{x}_m, \mathbf{x}_1) & k(\mathbf{x}_m, \mathbf{x}_2) & \dots & k(\mathbf{x}_m, \mathbf{x}_m) \end{bmatrix}$$

Kernel Gram Matrix

- It is not necessary to maintain all the original information

Kernels



- Three documents (d_1 , d_2 , and d_3) and 13 words in the dictionary ($t_1, t_2, \dots, t_{12}, t_{13}$)

term by document

$$D = [d_1, d_2, d_3] = \begin{matrix} & \begin{matrix} d_1 & d_2 & d_3 \end{matrix} \\ \begin{matrix} t_1 \\ t_2 \\ \dots \\ t_{13} \end{matrix} & \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 0 \\ \dots & \dots & \dots \\ 0 & 0 & 1 \end{bmatrix} \end{matrix}$$

doc by doc

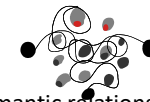
$$G = D'D = \begin{bmatrix} 0 & 1 & \dots & 0 \\ 1 & 1 & \dots & 0 \\ 0 & 0 & \dots & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 0 \\ \dots & \dots & \dots \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 4 & 3 & 0 \\ 3 & 5 & 1 \\ 0 & 1 & 5 \end{bmatrix}$$

K (Gram)

d_1 and d_2 share 3 words

- Size: 500 doc x 10.000 word \rightarrow 5.000.000 elements in matrix D
250.000 elements in matrix G

Kernels



- The inner product does not take into account the semantic relationship between terms:

$d_1 = \text{health, hospital, doctor} \dots$ $d_2 = \text{nurse, pills, } \dots$

$$k(d_1, d_2) = \langle \phi(d_1) \cdot \phi(d_2) \rangle = 0 \text{ identical words}$$

- It is possible to measure similarity between words instead of document similarity

$d_1 = x \rightarrow \phi(x) = \{0, 1, 1, 0, 0, \dots, 0, 1, 1, 0, 0, 0, 0\}$
 $d_2 = z \rightarrow \phi(z) = \{1, 1, 0, 0, 0, \dots, 0, 1, 1, 1, 0, 0, 0\}$
 $d_3 = v \rightarrow \phi(x) = \{0, 0, 0, 0, 0, \dots, 0, 0, 0, 1, 1, 1, 1\}$

term by term

$$DD' = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 0 \\ \dots & \dots & \dots \\ 1 & 1 & 0 \\ \dots & \dots & \dots \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 & 1 & \dots & 0 \\ 1 & 1 & \dots & 0 \\ 0 & 0 & \dots & 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 & \dots & 0 \\ 1 & 2 & \dots & 1 \\ \dots & \dots & \dots & \dots \\ 1 & 2 & \dots & 1 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & 1 \end{bmatrix} \quad T$$

0 docs with words t_1 and t_{13}

2 docs with words t_2 and t_7

Similarity between words. Different words but they appear together

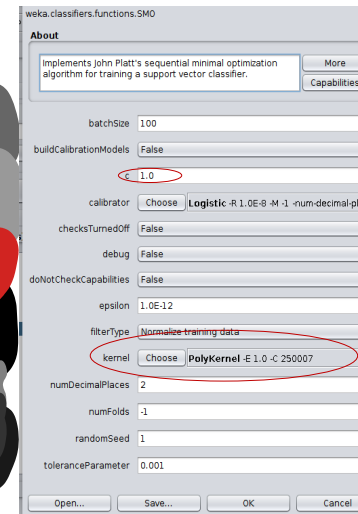
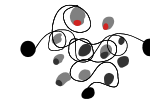
Kernels



- **Identity** (linear kernel) $K(\mathbf{x}, \mathbf{z}) = \langle \mathbf{x} \cdot \mathbf{z} \rangle$
 - **Polynomial** (polinomy of degree d) $K(\mathbf{x}, \mathbf{z}) = (\langle \mathbf{x} \cdot \mathbf{z} \rangle + c)^d$
- $$\langle \mathbf{x} \cdot \mathbf{z} \rangle^2 = \left(\sum_{i=1}^N x_i z_i \right)^2 = \left(\sum_{i=1}^N x_i z_i \right) \cdot \left(\sum_{j=1}^N x_j z_j \right) = \sum_{i=1}^N \sum_{j=1}^N x_i x_j z_i z_j = \sum_{i,j=1}^N (x_i x_j) (z_i z_j)$$
- **Gaussian** kernel (Radial Basis Function, RBF)
- $$K(\mathbf{x}, \mathbf{z}) = \exp\left(-\|\mathbf{x} - \mathbf{z}\|^2 / 2\sigma^2\right)$$
- **Sigmoid** $K(\mathbf{x}, \mathbf{z}) = \tanh(\kappa \langle \mathbf{x} \cdot \mathbf{z} \rangle + \vartheta)$

66

SMO (sequential minimal optimization)



Weka

- Normalizes attributes
- Converts nominal attributes to binary

Select C

- small C : great error tolerance → many training examples misclassified
- bigger C : results will improve
- very big C : great importance to the training data → overfitting

Select Kernel

- Polynomial (PolyKernel): exponent
- Gaussian kernel (RBF)

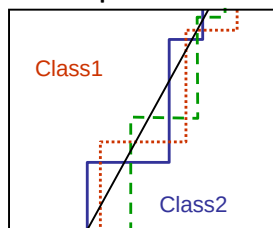
67

Multiple classifier systems

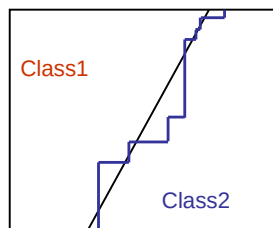


Motivation:

- **Statistical:** For small training samples the algorithm might find different hypothesis with the same performance. Taking into account several classifiers reduces the risk of selecting the wrong classifier.
- **Representational:** In many learning problems the objective function can not be represented for none of the classifiers.
- **DT example:**



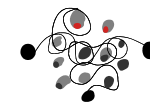
Decision Boundary (DB) of
3 individual DTs



Corresponding DB of
the voting classifier

68

Multiple classifier systems



T classifiers $\Phi_1, \Phi_2, \dots, \Phi_T$, to solve the same task

Classification based on **different learning methods**

result: combination of the result of k classifiers

- Voting in classification tasks
- Average result if it is numeric

T classifiers of the **same type**

Bagging: same classifier different subsamples. All classifiers same weight

Boosting: complementary classifiers. (sequential learning)

Random subspace methods: classifiers built with different sets of features

Weka: Vote (select T different classifiers)

AdaBoost, Bagging, Random subspace (select a classifier)

69

Multiple classifier systems: bagging



Bagging: Bootstrap aggregating, (Breiman en 1996)

Classifiers built using bootstrap samples (with replacement)

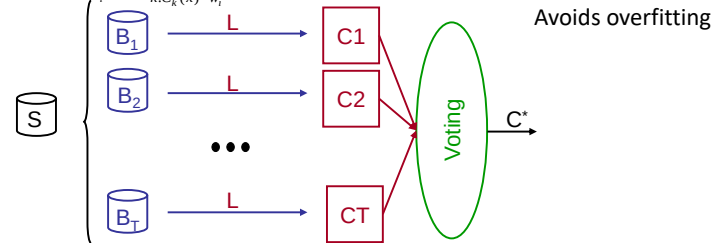
Generates **T subsamples** of size n' ($n' \leq n$) from S (learning sample of size n)

Builds **T classifiers** and combines the outputs

Final decision: **voting** of all individual classifiers for classification

Average for regression

$$C^*(x) = \arg \max_{w_i \in C} \sum_{k: C_k(x)=w_i} 1 \quad /* \text{The most voted class} */$$



70

Multiple classifier systems: boosting



Boosting: proposed with the aim of strengthening of weak classifiers (Schapire 1990)

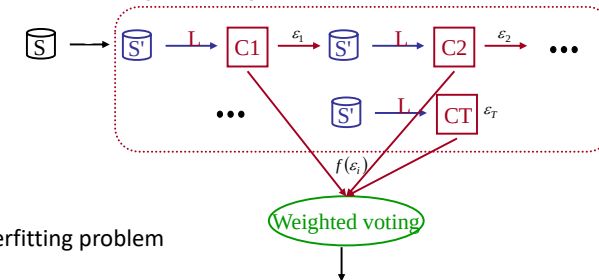
AdaBoost (Adaptive Boosting) introduced in by 1996 Freund&Schapire.

T classifiers are built sequentially. Each of the **instances** in the sample has a **weight** which changes depending on whether its classification is correct or incorrect.

resampling vs reweighting

- Weight increments for incorrectly classified examples
- Weight decrements for correctly classified examples

Final decision: weighted voting



71

Multiple classifier systems: RSM



Random subspace method (RSM) (Ho 1995: *Random Decision Forests*)

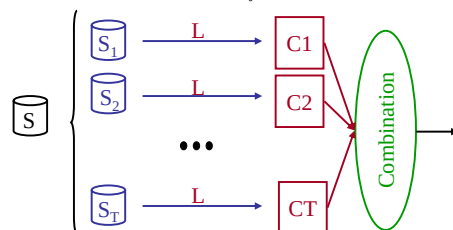
To build individual classifiers based on different subspaces

Subspace: classification space based on a subset of the original set of features

The methodology is applicable to any type of classifier

Final decision: average of class membership probabilities produced by each individual classifier.

$$C^*(x) = \arg \max_{w_i \in C} \frac{1}{T} \sum_{j=1}^T \hat{P}_j(w_i | x)$$



72

Example



• *soybean.arff* (%66)

- J48; RandomTree; IB1
- AdaBoost: J48, IB1, RTree;
- Bagging: J48, IB1, RTree;
- Vote (J48+RandomTree); (IB1+J48+RandomTree)

| | Basic | AdaBoost | Bagging | Vote |
|---------------|-------|----------|---------|------|
| IB1 | | | | |
| J48 | | | | |
| RandomTree | | | | |
| J48+RTree | | | | |
| IB1+J48+RTree | | | | |

73

Example



- *ReutersTrainGrain_WV.arff*, *soybean.arff*, *spam.arff* (%66)
 - IBK, Naive Bayes, J48, PART
 - F-measure/accuracy

| | lbk (k= ??) | Naive Bayes | J48 | PART |
|---------|-------------|-------------|-----|------|
| RGT_WV | F-m | F-m | F-m | F-m |
| | Acc | Acc | Acc | Acc |
| Soybean | F-m | F-m | F-m | F-m |
| | Acc | Acc | Acc | Acc |
| spam | F-m | F-m | F-m | F-m |
| | Acc | Acc | Acc | Acc |